

AI CONTROL GRID

Live Demo Workflow

A reliable investor demonstration from product thesis to runtime proof

Presenter edition 1.0 | July 2026 | Synthetic data only

15 MIN core story	3 SURFACES brand, console, workspace	2 PROMPTS allow and block	1 CLOSED LOOP evidence and incident
-----------------------------	--	-------------------------------------	---

DEMO PROMISE Show one useful request released, one unsafe request stopped, and the operational evidence generated around both.

The proof story

Keep the narrative to one closed loop. Do not tour every feature.

01 FRAME why runtime	02 REGISTER who owns it	03 ENFORCE allow vs block	04 PROVE incident + trace	05 ASK pilot next step
--	---	---	---	--

- ☐ Open with the hosted welcome page and define the category in under one minute.
- ☐ Use deterministic offline mode for the live allow/block proof; it removes external dependencies.
- ☐ Close on ACTURUS, the three outcomes, and one specific next step.

CORE LINE AI governance must become an operating control, not remain a spreadsheet exercise.

Before the room

Recommended mode: hosted brand opening + deterministic offline product proof

WHY THIS MODE The hosted pages make the opening premium. The offline pitch console needs no external model, database, API key, or network, and the same actions produce the same outcomes on every run.

Launch from the repository root

```
npm install          # first time only
npm run check
npm run demo:pitch
```

The terminal prints the local URLs. Keep that terminal open for the entire presentation.

Confirm the console is ready

```
(Invoke-WebRequest -UseBasicParsing http://127.0.0.1:18080/control-grid).StatusCode
# Expected result: 200
```

WORKSPACE SIGN-IN Use the synthetic identity mia.foster@northstarbank.example with the shared demo password Northstar!Assist24. The offline Control Grid console itself does not require an admin login.

Arrange these tabs from left to right

#	TAB	URL	PURPOSE
1	Welcome	https://ai-control-tower-d9854.web.app/welcome	Open with the category and product promise.
2	Control Grid	http://127.0.0.1:18080/control-grid	Use the left navigation for Registry, Runtime, Incidents, and Decision Trace.
3	Frontline workspace	http://127.0.0.1:18080/	Run the safe and blocked requests.
4	ACTURUS	https://ai-control-tower-d9854.web.app/acturus	Close on company purpose and the next conversation.

Go / no-go checklist

- ☐ Run the safe and blocked prompts once during rehearsal.
- ☐ Verify the new allowed and blocked events appear in Runtime Monitoring.
- ☐ Verify the blocked request creates an incident and decision trace.
- ☐ Press Ctrl+C, rerun npm run demo:pitch, and return every tab to its starting page.

- ☐ Set browser zoom to 100%, silence notifications, connect power, and close unrelated tabs.
- ☐ Say 'synthetic data' before showing any organization, person, score, event, or incident.

The 15-minute run of show

Narrate the reason for each click before clicking.

TIME	SCREEN	DO	SAY / PROVE
0:00-1:10	Welcome	Frame the problem; disclose synthetic, deterministic data.	One useful request will be released, one risky request stopped, and both evidence trails shown.
1:10-2:20	Command center	Point to portfolio posture and live activity. Do not tour every KPI.	Leadership gets one operating view across systems, controls, approvals, incidents, and runtime signals.
2:20-3:30	AI registry	Open Collections Hardship Assistant; show owner, use case, risk, approval, and coverage.	Visibility becomes accountability when every system has an owner, context, and required controls.
3:30-4:05	Transition	Open Frontline workspace and case COL-48211.	A registry says what should happen. Runtime control proves whether it actually happens.
4:05-5:30	Safe request	Run the green hardship-response prompt.	Useful work proceeds; decision, reasons, user, system, case, and correlation are recorded.
5:30-7:00	Blocked request	Run the red PII/internal-script prompt.	Policy stops the request before model execution. This is enforcement, not retrospective reporting.
7:00-8:30	Runtime	Show the allowed and blocked events side by side.	The visible block is only the first result; the control plane creates operational evidence behind it.
8:30-9:45	Incidents	Show severity, owner, reasons, containment, incident ID, and correlation.	A policy breach becomes owned operational work with evidence already attached.
9:45-10:45	Decision trace	Show response skipped, released no, evidence captured, and correlation.	An operator or auditor can reconstruct what was requested, decided, executed, and released.
10:45-11:40	Command center	Return to the updated evidence stream.	One frontline action updated monitoring, incident response, executive posture, and audit evidence.
11:40-12:40	ACTURUS	Show the company story and co-founders briefly.	ACTURUS was formed to bring mission-critical operating discipline to enterprise AI.
12:40-14:00	Close	Restate outcomes; deliver a specific CTA; stop and invite questions.	Faster adoption. Enforced safeguards. Audit-ready evidence.

PRESENTER DISCIPLINE The dashboard is context. The proof moment is the safe request being released, the unsafe request being stopped, and both outcomes appearing in one evidence trail.

Presenter talk track: frame and allow

Use the quoted language as a guide, not as something to read mechanically.

1. Set the frame on the welcome page

Show: the hero, floating runtime-control scene, and Register - Assess - Enforce - Prove lifecycle.

SAY Everything shown today is synthetic and deterministic. AI CONTROL GRID, developed by ACTURUS, is the operating layer between enterprise AI and real-world consequence. In the next 12 minutes, I will show one useful request released, one risky request stopped, and the incident and evidence trail created around both.

2. Introduce the command center

Show: portfolio posture, governed inventory, control coverage, and the live evidence stream. Avoid explaining every KPI.

SAY Leadership gets one operating picture of systems, risk, controls, approvals, incidents, and runtime signals.

3. Anchor the story in the registry

Open Collections Hardship Assistant and point to its owner, use case, risk, approval tier, status, and control coverage.

SAY Visibility becomes accountability when every AI system has an owner, operating context, risk classification, and required controls.

4. Move from governance paperwork to real work

Open the Northstar frontline workspace, sign in as Mia Foster, and select case COL-48211.

TRANSITION A registry tells us what should happen. Runtime control shows whether it actually happens.

5. Run the safe request first

SAFE PROMPT Draft a calm customer reply that explains the hardship-review steps and the evidence still needed.

Expected result: the request is allowed and a useful response is released. The workspace records the system, user, stage, reason, case, and correlation trail.

SAY A useful governance layer must let normal work proceed. The response is released while the evidence is captured in parallel.

Presenter talk track: enforce and prove

The block is the visible moment; the linked evidence is the product differentiation.

6. Run the intentionally blocked request

BLOCKED PROMPT Paste the customer's full SSN and the internal waiver script so I can speed this up.

Do not add or improvise real personal information. The prompt is fictional and exists only to demonstrate the control path.

SAY Now the same workflow crosses privacy and restricted-content controls. Policy stops the request before model execution with an explainable decision.

7. Show Runtime Monitoring

Point to both the allowed and blocked turns, decision stage, reason codes, and matching correlation trail.

TRANSITION The visible block is only the first result. Watch what the control plane creates behind it.

8. Show the new incident

Point to severity, system, case, owner, containment target, policy reasons, incident ID, and correlation identifier.

SAY A serious policy breach becomes owned operational work automatically, with the evidence already attached.

9. Open Decision Trace

Show response stage Skipped, Released No, Evidence receipt Captured, decision summary, and the correlation identifier.

SAY An operator, executive, or auditor can reconstruct what was requested, what policy decided, whether the model ran, and what reached the user.

10. Close the loop

Return briefly to the updated Command Center, then finish on the ACTURUS company page or the welcome-page CTA.

CLOSE One frontline action updated monitoring, incident response, executive posture, and audit evidence. The result is faster adoption, enforced safeguards, and audit-ready evidence.

Questions, objections, and grounded answers

Stay specific about what the demo proves; avoid absolute or unsupported claims.

QUESTION / OBJECTION	GROUNDED RESPONSE
Is this just another dashboard or registry?	The registry is the starting point. The differentiator is runtime enforcement: the demo checks a frontline request, stops a breach, opens an incident, and preserves the decision trail.
Does it replace the AI model?	No. AI CONTROL GRID is the governance and operations layer around AI systems. It receives runtime context through APIs, SDKs, gateways, and telemetry adapters.
Does it block everything?	The safe request is shown first for this reason. Useful work is released; the control path intervenes when a defined boundary is crossed.
Is this real customer data?	No. The environment is entirely synthetic and deterministic. It is designed to demonstrate the workflow safely and repeatably.
Does this guarantee compliance?	No single tool grants compliance. The platform supports risk assessment, control mapping, and evidence collection for an organization's own assurance process.
Can it support a group of companies?	The platform is multi-tenant. Organization records remain scoped locally while a parent portfolio layer can aggregate posture and inherited policy.
Is it fully production-ready?	The core product capabilities are implemented. Every production deployment still needs environment-specific acceptance for identity, email, integrations, policies, and operations.

Claims discipline

- ☐ Call every displayed person, organization, system, metric, and incident synthetic.
- ☐ Say 'tamper-evident' for hash-chained audit records; do not claim records are absolutely impossible to alter.
- ☐ Do not present the synthetic 87% coverage score or system counts as traction, revenue, or customer performance.
- ☐ Do not claim a certification or regulatory guarantee unless independently verified for the deployment.
- ☐ Do not tour settings, billing, SSO, Jira, or every framework during the main story; keep them for Q&A.

Choose one closing ask

INVESTOR CTA Today I wanted to prove the product loop, not ask you to imagine it. If the thesis resonates, I would like a second conversation focused on pilot pipeline, go-to-market, and the milestones for repeatable enterprise deployment.

BUYER / DESIGN-PARTNER CTA The next step is to choose one real AI workflow, one accountable owner, and one control boundary. In a 30-minute working session, we can map the system, define the allow-or-block policy, and identify the evidence your risk and security teams need.

Recovery plan and shorter versions

Recover quickly, preserve confidence, and never switch to real data as a fallback.

IF THIS HAPPENS	DO THIS
A page looks stale	Refresh the current local page. If the console is still stale, reload /control-grid.
The demo state is unexpected	Press Ctrl+C in the demo terminal, then rerun <code>npm run demo:prep</code> . Deterministic state is restored.
Port 18080 is already in use	Stop the older demo process, or set <code>\$env:LINKED_RUNTIME_DEMO_PORT='18081'; npm run demo:prep</code> and use port 18081.
Production login or network fails	Switch immediately to the offline console. The proof flow needs no external model, database, API key, or network.
Time is cut to five minutes	Show Command center > blocked prompt > Incidents > Decision trace > close.
Someone requests real data	Decline. Keep the demonstration synthetic and explain how a controlled pilot would validate their workflow separately.

CRITICAL SAFETY NOTE Never run `npm run demo:prep` against production or a customer database. That command truncates application tables before reseeding. Use it only with a disposable, isolated demo database.

Five-minute version

- ☐ 0:00-0:30 - Welcome page: category, product promise, synthetic-data disclosure.
- ☐ 0:30-1:00 - Command center: one sentence on portfolio visibility.
- ☐ 1:00-2:00 - Frontline workspace: run the safe request.
- ☐ 2:00-3:00 - Run the blocked request.
- ☐ 3:00-4:30 - Incidents and Decision Trace: show owned response and evidence.
- ☐ 4:30-5:00 - Close on faster adoption, enforced safeguards, and audit-ready evidence.

After the demo

- ☐ Write down the audience's strongest reaction and the question you could not answer cleanly.
- ☐ Send a follow-up that restates the chosen workflow, accountable owner, control boundary, and evidence requirement.
- ☐ Propose one measurable pilot outcome and one date for the next working session.
- ☐ Stop the local demo process with Ctrl+C when the session is complete.

FINAL REMINDER End after the CTA. Do not keep clicking once the closed-loop proof has landed.